

exercise_plotting_part1

October 8, 2025

1 Plotting Exercises, Part 1

1.0.1 Exercise 1

Create a pandas dataframe from the “Datasaurus.txt” file using the code:

```
[18]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

pd.set_option("mode.copy_on_write", True)

df = pd.read_csv(
    "https://raw.githubusercontent.com/nickeubank/practicaldatascience"
    "/master/Example_Data/Datasaurus.txt",
    delimiter="\t",
)
```

```
[19]: df
```

```
[19]:
```

	example1_x	example1_y	example2_x	example2_y	example3_x	example3_y	\
0	32.331110	61.411101	51.203891	83.339777	55.993030	79.277264	
1	53.421463	26.186880	58.974470	85.499818	50.032254	79.013071	
2	63.920202	30.832194	51.872073	85.829738	51.288459	82.435940	
3	70.289506	82.533649	48.179931	85.045117	51.170537	79.165294	
4	34.118830	45.734551	41.683200	84.017941	44.377915	78.164628	
..	
137	59.851838	72.958391	50.967748	29.679774	39.921363	19.701850	
138	48.960460	72.629526	91.191054	46.674343	84.794278	55.568650	
139	46.844855	36.791714	55.863768	85.336487	55.662959	83.356480	
140	39.963022	42.944915	49.280595	84.048823	50.492248	78.997532	
141	66.704944	32.015095	43.368502	84.332177	51.467101	79.201845	
	example4_x	example4_y	example5_x	example5_y	...	example9_x	\
0	55.3846	97.1795	51.147917	90.867412	...	47.695201	
1	51.5385	96.0256	50.517126	89.102395	...	44.609976	
2	46.1538	94.4872	50.207480	85.460047	...	43.856381	
3	42.8205	91.4103	50.069482	83.057670	...	41.578929	

4	40.7692	88.3333	50.562846	82.937822	...	49.177419
..
137	39.4872	25.3846	50.533635	17.019581	...	31.333244
138	91.2821	41.5385	77.500907	50.166986	...	86.401550
139	50.0000	95.7692	50.691124	87.513960	...	47.442112
140	47.9487	95.0000	49.990395	83.997357	...	46.264741
141	44.1026	92.6923	50.127182	82.990750	...	40.163816

	example9_y	example10_x	example10_y	example11_x	example11_y	\
0	95.241187	58.213608	91.881892	50.481508	93.222701	
1	93.075835	58.196054	92.214989	50.282406	97.609984	
2	94.085872	58.718231	90.310532	50.186703	99.694680	
3	90.303567	57.278373	89.907607	50.326911	90.022053	
4	96.610532	58.082020	92.008145	50.456207	89.987410	
..	
137	32.538569	43.722551	19.077328	30.487392	19.779470	
138	38.746933	79.326078	52.900391	89.500180	31.978917	
139	98.184302	56.663974	87.940125	50.410272	98.628369	
140	94.116192	57.821789	90.693167	50.325924	94.994631	
141	87.448672	58.243172	92.104328	50.104031	95.088538	

	example12_x	example12_y	example13_x	example13_y
0	65.815540	95.588374	38.337757	92.472719
1	65.672265	91.933402	35.751871	94.116768
2	39.002716	92.261838	32.767218	88.518295
3	37.795303	93.532455	33.729607	88.622266
4	35.513901	89.599190	37.238249	83.724928
..
137	33.674442	26.090490	34.794594	13.969683
138	75.627255	37.128752	79.221764	22.094591
139	40.610125	89.136240	36.030880	93.121733
140	39.114366	96.481751	34.499558	86.609985
141	34.583829	89.588902	31.106867	89.461635

[142 rows x 26 columns]

Note that the file being downloaded is *not* actually a CSV file. It is tab-delimited, meaning that within each row, columns are separated by tabs rather than commas. We communicate this to pandas with the `delimiter="\t"` option ("`\t`" is how we write a tab, as we will discuss in future lessons).

1.0.2 Exercise 2

This dataset actually contains 13 separate example datasets, each with two variables named `example[number]_x` and `example[number]_y`.

In order to get a better sense of what these datasets look like, write a loop that iterates over each example dataset (numbered 1 to 13) and print out the mean and standard deviation for `example[number]_x` and `example[number]_y` for each dataset.

For example, the first iteration of this loop might return something like:

```
Example Dataset 1:
Mean x: 23.12321978429576,
Mean y: 98.23980921730972,
Std Dev x: 21.2389710287,
Std Dev y: 32.2389081209832,
Correlation: 0.73892819281
```

(Though you shouldn't get those specific values. You might get values that are quite similar across datasets.)

Hint: When writing this type of code, it is often best to start by writing code to do what you want for the first iteration of the loop. Or, as Drew and Genevieve would say, **WORK ONE CASE BY HAND!** Once you have code that works for the first example dataset, then write the full loop around it.

```
[20]: for num in range(len(df.columns) // 2):
        mean_x = df[f"example{num + 1}_x"].mean()
        mean_y = df[f"example{num + 1}_y"].mean()
        std_x = df[f"example{num+1}_x"].std()
        std_y = df[f"example{num+1}_y"].std()
        corr = np.corrcoef(df[f"example{num + 1}_x"], df[f"example{num + 1}_y"])[0,1]
        print(
            f"Example Dataset {num+1}:\n mean x: {mean_x}\n mean y: {mean_y}\n
            standard deviation x: {std_x}\n standard deviation y: {std_y}\n Correlation:
            {corr}"
        )
```

```
Example Dataset 1:
mean x: 54.266099784295776
mean y: 47.834720624943664
standard deviation x: 16.769824954043756
standard deviation y: 26.9397434192671
Correlation: -0.0641283521673984
```

```
Example Dataset 2:
mean x: 54.268730022394365
mean y: 47.83082315530282
standard deviation x: 16.769239493454403
standard deviation y: 26.935726689918784
Correlation: -0.06858639424107654
```

```
Example Dataset 3:
mean x: 54.26731970598592
mean y: 47.83771726725352
standard deviation x: 16.76001265980608
standard deviation y: 26.930036087838204
Correlation: -0.06834335648025565
```

```
Example Dataset 4:
```

mean x: 54.26327323943662
mean y: 47.832252816901416
standard deviation x: 16.76514203911679
standard deviation y: 26.935403486939116
Correlation: -0.06447185270095167

Example Dataset 5:
mean x: 54.26030345169014
mean y: 47.839829209014084
standard deviation x: 16.767735488473807
standard deviation y: 26.93019151853346
Correlation: -0.06034144199921764

Example Dataset 6:
mean x: 54.26144178316902
mean y: 47.83025191366197
standard deviation x: 16.765897903899337
standard deviation y: 26.93987622043797
Correlation: -0.06171483797263011

Example Dataset 7:
mean x: 54.26880527950703
mean y: 47.83545020401409
standard deviation x: 16.766704015934764
standard deviation y: 26.939997961411027
Correlation: -0.06850422049412316

Example Dataset 8:
mean x: 54.26784882366197
mean y: 47.83589633112676
standard deviation x: 16.76675894771805
standard deviation y: 26.936104931679978
Correlation: -0.068979735359512

Example Dataset 9:
mean x: 54.26588178542254
mean y: 47.831495652323945
standard deviation x: 16.768852670828494
standard deviation y: 26.93860807087184
Correlation: -0.06860920641825635

Example Dataset 10:
mean x: 54.26734110478873
mean y: 47.83954522535211
standard deviation x: 16.76895921619445
standard deviation y: 26.93027468808843
Correlation: -0.0629611002206542

Example Dataset 11:
mean x: 54.26992723091549
mean y: 47.836987988408445
standard deviation x: 16.769958611325382
standard deviation y: 26.937683806980512
Correlation: -0.06944556959350369

Example Dataset 12:

```
mean x: 54.266916301197185
mean y: 47.83160198797184
standard deviation x: 16.769999617573024
standard deviation y: 26.937901927731797
Correlation: -0.06657523020460904
Example Dataset 13:
mean x: 54.26015033415493
mean y: 47.839717279450696
standard deviation x: 16.76995769550748
standard deviation y: 26.93000168716234
Correlation: -0.06558333729297582
```

1.0.3 Exercise 3



Based only on these results, discuss what might you conclude about these example datasets with your partner. Write down your thoughts.


Based on this data, the distribution of values between column x and y for each pair of columns is very similar, seeing as the mean and standard deviation values along with the correlation coefficient are approximately the same.



1.0.4 Exercise 4


Write a loop that iterates over these example datasets and plot a simple scatter plot of each dataset with the x variable on the x-axis and the y variable on the y-axis.



```
[21]: for num in range(len(df.columns) // 2):
      x = df[f"example{num + 1}_x"]
      y = df[f"example{num + 1}_y"]
      fig, ax = plt.subplots()
      ax.scatter(x, y)
      plt.show()
```
























1.0.5 Exercise 5

Review your plots. How does your impression of how these datasets differ from what you wrote down in Exercise 3?

We are thoroughly surprised by how differently scattered these points are given the summary statistics were approximately the same. We initially expected the graphs to look similar due to this. However, after giving it some thought, it seems reasonable that different distributions can give similar summary statistics since they just need to have a similar center and spread of data, even if the data points might be differently scattered.

1.1 Economic Development and... Your Choice!

1.1.1 Exercise 6

Load the World Development Indicator data [here](#)

Rather than picking a single year, pick a single country and look at how GDP per capita and one of the other variables in that dataset have evolved together over time.

Make any adjustments to the functional forms of your variables and/or axes needed to make the figure legible.

```
[25]: world_df = pd.read_csv(
        "http://raw.githubusercontent.com/nickeubank/practicaldatascience/master/
        ↪Example_Data/wdi_plotting.csv"
    )
    preferred_country = "Argentina"
    country_df = world_df[world_df["Country Name"] == preferred_country].copy()

[ ]: country_df = country_df.sort_values("Year")
    gdp_col = "GDP per capita (constant 2010 US$)"
    life_col = "Life expectancy at birth, total (years)"
    country_df = country_df[["Year", gdp_col, life_col]].dropna()
    country_df["log_GDP"] = np.log10(
        country_df[gdp_col]
    ) # log transform for easier understanding
```

1.1.2 Exercise 7

Now add a second series. Create a pair of plots so that the two subplots are positioned so that they are effectively sharing the same time axes (e.g., if you draw a line up from 2010 on one plot, you get to 2010 on the other).


Use your detective skills (and some guess and check work) to figure out how to get it to work!

```
[ ]: # gdp plot - top
    fig, (ax1, ax2) = plt.subplots(2, 1, figsize=(10, 6), sharex=True)
    ax1.plot(country_df["Year"], country_df["log_GDP"], marker="o", color="blue")
    ax1.set_ylabel("log (GDP per capita)")
    ax1.set_title(f"{preferred_country}: GDP per capita and Life Expectancy Over_
    ↪Time")

    # life expectancy plot - bottom
    ax2.plot(country_df["Year"], country_df[life_col], marker="o", color="green")
    ax2.set_ylabel("Life Expectancy (years)")
    ax2.set_xlabel("Year")

    plt.tight_layout()
    plt.show()

    corr = np.corrcoef(country_df["log_GDP"], country_df[life_col])[0, 1]
    print(f"Correlation between log(GDP per capita) and life expectancy: {corr:.
    ↪3f}")
```



Correlation between $\log(\text{GDP per capita})$ and life expectancy: 0.737

- From 1970 to 2020, Argentina's GDP per capita goes up and down quite a bit instead of following a smooth upward path. There are clear dips in the late 1980s, early 2000s, and again around 2018–2020, which line up with known economic crises in the country. Even with these ups and downs, GDP per capita still trends slightly upward overall across the 50 years.
- On the other hand, life expectancy (the bottom graph) shows a steady and consistent increase from about 67 years in 1970 to over 76 years in 2020. This means that people in Argentina have been living longer over time, likely because of improvements in healthcare, living conditions, and technology, even when the economy was struggling.

-Overall, there seems to be a positive relationship between GDP per capita and life expectancy in the long run — as the economy grows, people tend to live longer. But the steady rise in life expectancy even when GDP falls suggests that factors other than income, like medical advancements and public health programs, also play a big role in improving overall well-being.